

Homework 4

Shun-Lung Chang, Dilip Hiremath

```
library(foreign)
hsb <- read.dta('data/hsbdemo.dta')
```

1. Cross-tabulate the variables ses and prog.

```
ct <- table(hsb$prog, hsb$ses) / nrow(hsb)
ct
```

	low	middle	high
general	0.080	0.100	0.045
academic	0.095	0.220	0.210
vocation	0.060	0.155	0.035

(a) (half a point) Which program was chosen by the largest fraction of students with high socio-economic status?

(b) (half a point) How many percent of students with low socio-economic status selected the general program?

(c) (half a point) In the academic program are there more students with middle socioeconomic status than students with high socio-economic status?


(d) (half a point) What is the least-frequent combination of the two variables?

2. You continue with your analysis of the relationship between ses and prog.

(a) (half a point) Draw a mosaicplot visualising the contingency table of program choice and socio-economic status.

```
mosaicplot(ct, main = 'Program vs. Socio-economic Status')
```

Program vs. Socio-economic Status



(b) (1.5 points) Are students with low ses less likely (as measured in odds) to choose the academic program than students with higher socio-economic status? Calculate the odds ratios for choosing the academic program comparing students with low ses to students with middle ses and to students with high ses. [hint: use the command `loddsratio` in the package `vcd`. First, aggregate the variable `prog` into a binary variable indicating whether the student has chosen an academic program yes or no.]

```
library(vcd)
hsb$academic_yes <- hsb$prog == 'academic'

ct_2 <- table(hsb$academic_yes, hsb$ses) / nrow(hsb)
loddsratio(ct_2)
```

log odds ratios for and

```
low:middle middle:high
0.2401295 1.1127169
```

3. Now, you assess the relationship between `prog` and `ses` using the χ^2 -statistic.

(a) (1 point) Calculate the χ^2 -test to assess the relationship between `ses` and `prog`. Is the relationship statistically significant?

```
chisq.test(hsb$prog, hsb$ses)
```

Pearson's Chi-squared test

```
data: hsb$prog and hsb$ses
X-squared = 16.604, df = 4, p-value = 0.002307
```

(b) (1 point) Calculate the expected frequencies under the assumption that socio-economic status has no effect on program choice. For which cells are expected frequencies higher than the observed ones?

```
expected <- chisq.test(hsb$prog, hsb$ses)$expected
expected - (ct * nrow(hsb))
```

	hsb\$ses		
hsb\$prog	low	middle	high
general	-5.425	1.375	4.050
academic	5.675	5.875	-11.550
vocation	-0.250	-7.250	7.500

4. In the following, perform the last analysis separately for female and male students.

```
hsb_m <- hsb[hsb$female == 'male', ]
hsb_f <- hsb[hsb$female == 'female', ]
```

(a) (half a point) Calculate the χ^2 -test to assess the relationship between ses and prog.

```
chisq.test(hsb_m$prog, hsb_m$ses)
```

Pearson's Chi-squared test

```
data: hsb_m$prog and hsb_m$ses
X-squared = 11.557, df = 4, p-value = 0.02097
```

```
chisq.test(hsb_f$prog, hsb_f$ses)
```

Pearson's Chi-squared test

```
data: hsb_f$prog and hsb_f$ses
X-squared = 7.473, df = 4, p-value = 0.1129
```

(b) (half a point) Calculate the expected frequencies under the assumption that socio-economic status has no effect on program choice. For which cells are expected frequencies higher than the observed ones?

```
# male
observed_m <- table(hsb_m$prog, hsb_m$ses)
chisq.test(hsb_m$prog, hsb_m$ses)$expected - observed_m
```

	hsb_m\$ses		
hsb_m\$prog	low	middle	high
general	-3.5384615	0.8461538	2.6923077
academic	3.7472527	2.2747253	-6.0219780
vocation	-0.2087912	-3.1208791	3.3296703

```
# female
observed_f <- table(hsb_f$prog, hsb_f$ses)
chisq.test(hsb_f$prog, hsb_f$ses)$expected - observed_f
```


```
      hsb_f$ses
hsb_f$prog  low  middle  high
general -1.9541284 0.5688073 1.3853211
academic 2.0275229 3.5412844 -5.5688073
vocation -0.0733945 -4.1100917 4.1834862
```

(c) (half a point) Do the results differ for the two sexes?

(d) (half a point) Visualise the relationships using mosaicplots. Get any differences between females and males in relation to socio-economic status and program choice visible in the plots?


```
mosaicplot(observed_m, 'Program vs. Socio-economic Status (Male)')
```

Program vs. Socio-economic Status (Male)



```
mosaicplot(observed_f, 'Program vs. Socio-economic Status (Female)')
```

Program vs. Socio-economic Status (Female)



5. Create a multinomial logistic regression model using prog as dependent variable and the following predictors: female, ses, schtype, read, write, math, science, honors, awards. [hint: use the function multinom in the package nnet.]

```
library(nnet)
mod_1 <- multinom(prog ~ . - id - socst - cid,
                  data = hsb)
```

```
# weights: 36 (22 variable)
initial value 219.722458
iter 10 value 184.828226
iter 20 value 157.499339
final value 157.443537
converged
```

```
summary(mod_1)
```

Call:

```
multinom(formula = prog ~ . - id - socst - cid, data = hsb)
```

Coefficients:

	(Intercept)	female	female	sesmiddle	seshigh	schtypeprivate
academic	-5.498379	-0.1522507	0.3751645	1.0802062	0.5372474	
vocation	4.163618	0.2804404	1.3235922	0.7959664	-1.2758046	

	read	write	math	science	honors	enrolled
academic	0.052955458	0.06429399	0.10006643	-0.10199019	0.5743597	
vocation	-0.009609139	-0.02166161	-0.02995725	-0.03800852	1.9991952	

	awards
academic	-0.2703505
vocation	-0.3525181

Std. Errors:

	(Intercept)	female	female	sesmiddle	seshigh	schtypeprivate
academic	2.347466	0.4533749	0.5061641	0.5772063	0.5550686	
vocation	2.455693	0.5179411	0.5388177	0.7084789	0.8817585	

	read	write	math	science	honors	enrolled
academic	0.02919946	0.04927104	0.03480492	0.03090596	0.8644712	
vocation	0.03322165	0.05054331	0.03756352	0.03224286	1.1038196	

	awards
academic	0.2936491
vocation	0.3856715

Residual Deviance: 314.8871

AIC: 358.8871

(a) (half a point) How large is the AIC score for this model?

The AIC score is 358.8871.

```
AIC(mod_1)
```

```
[1] 358.8871
```

(b) (1.5 points) The default output does not include p-values. Compute p-values based on the Wald-test statistics and determine the coefficients that are statistically significantly different from zero!

As the table below shows, the variables *academic:(Intercept)*, *academic:math*, *academic:science* and *vocation:sesmiddle* are significant under significance level of 5%.

```
# compute p-values by definition
# z <- summary(mod_1, Wald.ratios = TRUE)$Wald.ratios
# p <- (1 - pnorm(abs(z))) * 2
# p

# compute p-values by package function
library(AER)
coeftest(mod_1)[, 4] < 0.05
```

academic:(Intercept)	academic:femalefemale	academic:sesmiddle
TRUE	FALSE	FALSE
academic:seshigh	academic:schtypprivate	academic:read
FALSE	FALSE	FALSE
academic:write	academic:math	academic:science
FALSE	TRUE	TRUE
academic:honorsenrolled	academic:awards	vocation:(Intercept)
FALSE	FALSE	FALSE
vocation:femalefemale	vocation:sesmiddle	vocation:seshigh
FALSE	TRUE	FALSE
vocation:schtypprivate	vocation:read	vocation:write
FALSE	FALSE	FALSE
vocation:math	vocation:science	vocation:honorsenrolled
FALSE	FALSE	FALSE
vocation:awards		
FALSE		

6. Using the model from the previous question and the backward strategy with criterion AIC for variable selection, determine the significant coefficients in the resulting model.

```
mod_2 <- step(mod_1, direction = 'backward')
```

```
summary(mod_2)
```

Call:

```
multinom(formula = prog ~ ses + schtyp + read + math + science,
  data = hsb)
```

Coefficients:

	(Intercept)	sesmiddle	seshigh	schtypprivate	read
academic	-3.745688	0.323115	1.0358034	0.608257	0.05912408
vocation	3.907946	1.183126	0.7014962	-1.408038	-0.01218565
	math	science			
academic	0.10745053	-0.09076914			
vocation	-0.03078266	-0.04770648			

Std. Errors:

	(Intercept)	sesmiddle	seshigh	schtypprivate	read
academic	1.401302	0.4876102	0.5648570	0.5484718	0.02807034
vocation	1.564001	0.5201513	0.6880376	0.8662165	0.03127526

	math	science
academic	0.03270710	0.02859475
vocation	0.03535214	0.03005018

Residual Deviance: 322.9919

AIC: 350.9919

(a) (1 point) Which predictors are included in the resulting model?

According to the table above, the variables *ses*, *schtyp*, *read*, *math* and *science* are included in the final model.

(b) (half a point) What is the BIC score of the resulting model?

The BIC score is 397.1684.

```
BIC(mod_2)
```

```
[1] 397.1684
```

(c) (half a point) What is the log-likelihood score of this model?

The log-likelihood score is -161.496.

```
logLik(mod_2)
```

```
'log Lik.' -161.496 (df=14)
```

7. (2 points) Using the final model that resulted in Question 6 predict the probabilities for the three program types for the combination of all factor levels and the average score of numeric predictors in the model.

The table below indicates academic is the most possible outcome in all combinations. But the students in the private school are more likely to take an academic program than that in the public school.

```
d <- expand.grid(ses = c('low', 'middle', 'high'),
               schtyp = c('private', 'public'),
               read = mean(hsb$read),
               math = mean(hsb$math),
               science = mean(hsb$science))
```

```
pred <- predict(mod_2, newdata = d, type = 'probs', se = TRUE)
cbind(d, pred)
```

	ses	schtyp	read	math	science	general	academic	vocation
1	low	private	52.23	52.645	51.85	0.2801551	0.6897408	0.03010409
2	middle	private	52.23	52.645	51.85	0.2104441	0.7157334	0.07382253
3	high	private	52.23	52.645	51.85	0.1226533	0.8507663	0.02658044
4	low	public	52.23	52.645	51.85	0.3597989	0.4821528	0.15804831
5	middle	public	52.23	52.645	51.85	0.2333606	0.4319957	0.33464371
6	high	public	52.23	52.645	51.85	0.1766363	0.6668812	0.15648254

8. (2 points) Again using the final model that resulted in Question 6, we now want to investigate the specific dependency on the math score. Generate new data such that you have for each combination of factor levels a total of 51 math scores running from 30 to 80 in increments of one. The other numeric predictors enter again with their mean score into the prediction. Compute the predictions and average them for each level of socio-economic status.

The table below shows the average probabilities of taking different programs with respect to different socio-economic levels. As can be seen from the table, taking an academic program is the most probable in three socio-economic levels. Furthermore, people with a high socio-economic status is more likely to have an academic program.

```
d <- expand.grid(ses = c('low', 'middle', 'high'),
               schtyp = c('private', 'public'),
               read = mean(hsb$read),
               math = 30:80,
               science = mean(hsb$science))

pred <- predict(mod_2, newdata = d, type = 'probs', se = TRUE)

bind_d <- cbind(d, pred)
aggregate(bind_d[, 6:8], by = list(bind_d$ses), FUN = 'mean')
```

	Group.1	general	academic	vocation
1	low	0.2978392	0.5956783	0.1064825
2	middle	0.1993257	0.5887810	0.2118934
3	high	0.1629723	0.7138563	0.1231714